

## Neural Networks image recognition - ConvNet

1. Add random noise (see below on `size` parameter on `np.random.normal` (<https://numpy.org/doc/stable/reference/random/generated/numpy.random.normal.html>)) to the images in training and testing. \*\*Make sure each image gets a different noise feature added to it. Inspect by printing out several images. Note - the `size` parameter should match the data. \*\*
2. Compare the `accuracy` of train and val after N epochs for MLNN with and without noise.
3. Vary the amount of noise by changing the `scale` parameter in `np.random.normal` by a factor. Use `.1`, `.5`, `1.0`, `2.0`, `4.0` for the `scale` and keep track of the `accuracy` for training and validation and plot these results.
4. Compare these results with the previous week where we used a MultiLayer Perceptron (this week we use a ConvNet).

## Neural Networks - Image Recognition

```
In [1]: import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.optimizers import RMSprop
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend

import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

WARNING:tensorflow:From C:\Users\Wkcosm\Anaconda3\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse\_softmax\_cross\_entropy is deprecated. Please use tf.compat.v1.losses.sparse\_softmax\_cross\_entropy instead.

### Conv Net

Trains a simple convnet on the MNIST dataset. Gets to 99.25% test accuracy after 12 epochs (there is still a lot of margin for parameter tuning).

```
In [2]: # input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()

if backend.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

## 1. Add noise

```
In [3]: # Noise is added here
# The max value of the noise should not grossly surpass 1.0
noise_train = np.random.normal(0.5, 0.05, (60000, 28, 28, 1))
noise_test = np.random.normal(0.5, 0.05, (10000, 28, 28, 1))
x_train2 = x_train + noise_train
x_test2 = x_test + noise_test
```

```
In [4]: #check the minimum and maximum value of the noise
```

```
print(np.max(noise_train))
print(np.min(noise_train))
print(np.max(noise_test))
print(np.min(noise_test))
```

```
0.7937039777445263
0.22592893773982398
0.790001239714075
0.23102239788898254
```

## 2. Compare Accuracy

### 2-A. Accuracy without noise = 83.89%

```
In [5]: batch_size = 128
num_classes = 10
epochs = 12

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)

model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))

model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

model.fit(x_train, y_train,
          batch_size=batch_size,
          epochs=epochs,
          verbose=1,
          validation_data=(x_test, y_test))
score = model.evaluate(x_test, y_test, verbose=0)

print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

WARNING:tensorflow:From C:\Users\Wkcosm\Anaconda3\Lib\site-packages\keras\src\backend.py:873: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From C:\Users\Wkcosm\Anaconda3\Lib\site-packages\keras\src\layers\pooling\max\_pooling2d.py:161: The name tf.nn.max\_pool is deprecated. Please use tf.nn.max\_pool2d instead.

Epoch 1/12

WARNING:tensorflow:From C:\Users\Wkcosm\Anaconda3\Lib\site-packages\keras\src\utils\tf\_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\Wkcosm\Anaconda3\Lib\site-packages\keras\src\engine\base\_layer\_utils.py:384: The name tf.executing\_eagerly\_outside\_functions is deprecated. Please use tf.compat.v1.executing\_eagerly\_outside\_functions instead.

469/469 [=====] - 67s 137ms/step - loss: 2.2815 - accuracy: 0.1424 - val\_loss: 2.2540 - val\_accuracy: 0.2601

Epoch 2/12

469/469 [=====] - 71s 150ms/step - loss: 2.2349 - accuracy: 0.2375 - val\_loss: 2.1969 - val\_accuracy: 0.4094

Epoch 3/12

469/469 [=====] - 64s 136ms/step - loss: 2.1766 - accuracy: 0.3216 - val\_loss: 2.1215 - val\_accuracy: 0.5204

Epoch 4/12

469/469 [=====] - 63s 134ms/step - loss: 2.0958 - accuracy: 0.4020 - val\_loss: 2.0188 - val\_accuracy: 0.6073

Epoch 5/12

469/469 [=====] - 68s 144ms/step - loss: 1.9894 - accuracy: 0.4690 - val\_loss: 1.8816 - val\_accuracy: 0.6847

Epoch 6/12

469/469 [=====] - 65s 139ms/step - loss: 1.8534 - accuracy: 0.5320 - val\_loss: 1.7092 - val\_accuracy: 0.7397

Epoch 7/12

469/469 [=====] - 58s 123ms/step - loss: 1.6922 - accuracy: 0.5810 - val\_loss: 1.5123 - val\_accuracy: 0.7718

Epoch 8/12

469/469 [=====] - 59s 127ms/step - loss: 1.5241 - accuracy: 0.6144 - val\_loss: 1.3116 - val\_accuracy: 0.7962

Epoch 9/12

469/469 [=====] - 68s 145ms/step - loss: 1.3667 - accuracy: 0.6445 - val\_loss: 1.1293 - val\_accuracy: 0.8099

Epoch 10/12

469/469 [=====] - 60s 127ms/step - loss: 1.2290 - accuracy: 0.6694 - val\_loss: 0.9780 - val\_accuracy: 0.8234

Epoch 11/12

469/469 [=====] - 63s 135ms/step - loss: 1.1168 - accuracy: 0.6922 - val\_loss: 0.8607 - val\_accuracy: 0.8319

Epoch 12/12

469/469 [=====] - 58s 124ms/step - loss: 1.0278 - accuracy: 0.7056 - val\_loss: 0.7708 - val\_accuracy: 0.8389

Test loss: 0.7707778215408325

Test accuracy: 0.8389000296592712

**2-B. Accuracy with noise = 87.22%**

```
In [6]: model.fit(x_train2, y_train,
                  batch_size=batch_size,
                  epochs=epochs,
                  verbose=1,
                  validation_data=(x_test2, y_test))
score = model.evaluate(x_test2, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
Epoch 1/12
469/469 [=====] - 62s 131ms/step - loss: 1.2414 - accuracy: 0.5860 - val_loss: 0.7409 - val_accuracy: 0.8406
Epoch 2/12
469/469 [=====] - 58s 125ms/step - loss: 1.0687 - accuracy: 0.6628 - val_loss: 0.7181 - val_accuracy: 0.8469
Epoch 3/12
469/469 [=====] - 58s 125ms/step - loss: 1.0200 - accuracy: 0.6833 - val_loss: 0.6873 - val_accuracy: 0.8491
Epoch 4/12
469/469 [=====] - 57s 121ms/step - loss: 0.9890 - accuracy: 0.6926 - val_loss: 0.6576 - val_accuracy: 0.8518
Epoch 5/12
469/469 [=====] - 58s 125ms/step - loss: 0.9613 - accuracy: 0.7024 - val_loss: 0.6308 - val_accuracy: 0.8561
Epoch 6/12
469/469 [=====] - 56s 119ms/step - loss: 0.9321 - accuracy: 0.7112 - val_loss: 0.6063 - val_accuracy: 0.8580
Epoch 7/12
469/469 [=====] - 58s 123ms/step - loss: 0.9051 - accuracy: 0.7191 - val_loss: 0.5850 - val_accuracy: 0.8608
Epoch 8/12
469/469 [=====] - 58s 124ms/step - loss: 0.8856 - accuracy: 0.7260 - val_loss: 0.5658 - val_accuracy: 0.8641
Epoch 9/12
469/469 [=====] - 59s 125ms/step - loss: 0.8629 - accuracy: 0.7293 - val_loss: 0.5488 - val_accuracy: 0.8674
Epoch 10/12
469/469 [=====] - 57s 122ms/step - loss: 0.8396 - accuracy: 0.7377 - val_loss: 0.5330 - val_accuracy: 0.8686
Epoch 11/12
469/469 [=====] - 64s 137ms/step - loss: 0.8250 - accuracy: 0.7397 - val_loss: 0.5189 - val_accuracy: 0.8709
Epoch 12/12
469/469 [=====] - 73s 155ms/step - loss: 0.8081 - accuracy: 0.7473 - val_loss: 0.5058 - val_accuracy: 0.8722
Test loss: 0.50580894947052
Test accuracy: 0.8722000122070312
```

**Result: Accuracy increased by 3.33%p due to the noise (Why??)****3. Vary noises**

```
In [7]: accuracy_score = []
noise_scale = [.1, .5, 1.0, 2.0, 4.0]

#set for loop to see scores for each scale of noise
for scale in noise_scale:

    noise_train = np.random.normal(0.5, scale, (60000, 28, 28, 1))
    noise_test = np.random.normal(0.5, scale, (10000, 28, 28, 1))
    x_train3 = x_train + noise_train
    x_test3 = x_test + noise_test

    history3 = model.fit(x_train3, y_train,
                        batch_size=batch_size,
                        epochs=epochs,
                        verbose=1,
                        validation_data=(x_test3, y_test))

    score3 = model.evaluate(x_test3, y_test, verbose=0)
    accuracy = score3[1]
    accuracy_score.append(accuracy)
```

469/469 [=====] - 70s 149ms/step - loss: 1.7834 - accuracy: 0.3857 - val\_loss: 1.6175 - val\_accuracy: 0.4648

Epoch 5/12

469/469 [=====] - 73s 156ms/step - loss: 1.7951 - accuracy: 0.3807 - val\_loss: 1.6242 - val\_accuracy: 0.4608

Epoch 6/12

469/469 [=====] - 75s 148ms/step - loss: 1.8616 - accuracy: 0.3892 - val\_loss: 1.5972 - val\_accuracy: 0.4643

Epoch 7/12

469/469 [=====] - 73s 155ms/step - loss: 1.1165 - accuracy: 0.6286 - val\_loss: 0.8411 - val\_accuracy: 0.7312

Epoch 8/12

469/469 [=====] - 77s 164ms/step - loss: 1.1225 - accuracy: 0.6255 - val\_loss: 0.8444 - val\_accuracy: 0.7299

Epoch 9/12

469/469 [=====] - 70s 149ms/step - loss: 1.1251 - accuracy: 0.6277 - val\_loss: 0.8462 - val\_accuracy: 0.7295

Epoch 10/12

469/469 [=====] - 70s 149ms/step - loss: 1.1251 - accuracy: 0.6277 - val\_loss: 0.8462 - val\_accuracy: 0.7295

Epoch 11/12

469/469 [=====] - 77s 164ms/step - loss: 1.1225 - accuracy: 0.6255 - val\_loss: 0.8444 - val\_accuracy: 0.7299

Epoch 12/12

469/469 [=====] - 73s 155ms/step - loss: 1.1165 - accuracy: 0.6286 - val\_loss: 0.8411 - val\_accuracy: 0.7312

Epoch 1/12

469/469 [=====] - 75s 148ms/step - loss: 1.8616 - accuracy: 0.3892 - val\_loss: 1.5972 - val\_accuracy: 0.4643

Epoch 2/12

469/469 [=====] - 73s 156ms/step - loss: 1.7951 - accuracy: 0.3807 - val\_loss: 1.6242 - val\_accuracy: 0.4608

Epoch 3/12

469/469 [=====] - 68s 144ms/step - loss: 1.7910 - accuracy: 0.3816 - val\_loss: 1.6225 - val\_accuracy: 0.4639


Epoch 4/12

469/469 [=====] - 70s 149ms/step - loss: 1.7834 - accuracy: 0.3857 - val\_loss: 1.6175 - val\_accuracy: 0.4648

```
In [8]: print(accuracy_score)

[0.8865000009536743, 0.8525999784469604, 0.7311999797821045, 0.4729999899864197, 0.2337000072002411]
```

```
In [9]: plt.scatter(noise_scale, accuracy_score)
plt.xlabel('Scale of Noise')
plt.ylabel('Accuracy')
plt.show()
```



**Result: The accuracy significantly dropped as scale of noise increased**

#### 4. Compare noises with the result of the previous week

Scale of Noise : 0.1                      0.5                      1.0                      2.0

4.0

- Accuracy Score Using Multilayer: [0.9778000116348267, 0.9319999814033508, 0.7793999910354614, 0.44780001044273376, 0.22589999437332153]
- Accuracy Score Using ConvNet: [0.8865000009536743, 0.8525999784469604, 0.7311999797821045, 0.4729999899864197, 0.2337000072002411]

**Result: In lower noises(scale  $\leq 1.0$ ), the accuracy scores of Multilayer were higher; in high noises(scale  $\geq 2.0$ ), the accuracy scores of ConvNet were higher.**